

AI Product Management

The Complete Handbook

From Foundations to Real-World Execution



TABLE OF CONTENT

Part I. Introduction: The End of the AI Winter and the Beginning of a New Product Era

1. The Market Dynamics of Generative AI

1.1 The Post-2022 Economic Shift: Valuation Explosions

- 1.1.1 Case Study: OpenAI, Anthropic, and XAI
- 1.1.2 Legacy Pivot: Microsoft, Google, and Meta's Trillion-Dollar Growth
- 1.1.3 The Hardware Monopoly: NVIDIA's 10x Transformation

1.2 The "Bubble" Test: Distinguishing Hype from Utility

- 1.2.1 Lessons from Blockchain, Crypto, and NFTs
- 1.2.2 The "Daily Use" Verdict: Integrating AI into Personal Life

2. Defining the AI Product Manager (AIPM)

2.1 The Broad Classification: AI-Enabled vs. AI Product PMs

2.2 The Spectrum of AI Product PM Roles

- 2.2.1 Core AI PMs: Infrastructure, Vector Databases, and Model Development
- 2.2.2 Applied AI PMs: Solving Customer Pain Points with AI Capabilities

2.3 Case Study Contrast: Building Gamin (Model) vs. Chat GPT (Application)

Part II: Technical Foundations—How AI Thinks

3. Large Language Model (LLM) Fundamentals

3.1 The Architecture: The "Next Token Predictor."

3.2 Tokenization: Numerical Language for Machines

- 3.2.1 Vocabulary Size and Memory Efficiency
- 3.2.2 Economics of Tokens: Input vs. Output Pricing

3.3 The Three Stages of LLM Creation

- 3.3.1 Pre-training: Crawling the Internet for Pattern Recognition
- 3.3.2 Training: Solving the Billion-Parameter Mathematical Equation
- 3.3.3 Post-training: Conditioning the "Helpful Assistant."

4. Advanced Neural Architectures

- 4.1 The Transformer Revolution (2017): "Attention is All You Need"
- 4.2 The Mechanism of Attention: Parallel Processing and Relationship Mapping
- 4.3 Hardware Requirements: Why LLMs Run on GPUs instead of CPUs
- 4.4 Stochastic Nature: Probabilistic vs. Deterministic Outputs

5. The Supervised Learning Framework

- 5.1 Human Labeling and Supervised Learning
- 5.2 RLHF: Reinforcement Learning from Human Feedback
- 5.3 Small Language Models (SLM) vs. Large Language Models (LLM)

Part III: The AI Product Development Life Cycle (AIDLC)

6. Phase 1: Product Discovery (The Problem Space)

- 6.1 Sources of Insights: Business Goals, Stakeholders, and Data Analytics
- 6.2 The First Principle: Starting from the Problem, Not the Solution

7. Phase 2: Solution Mapping (The 5 AI Capabilities)

- 7.1 Recommendations and Ranking (E-commerce and Streaming)
- 7.2 Categorization and Classification (Email and Spam)
- 7.3 Forecasting and Predictive Analytics (Sales and Weather)
- 7.4 Anomaly Detection (Fraud and Security)
- 7.5 Contextual Content Generation (The Generative AI Layer)

8. Phase 3: Product Delivery (Execution)

- 8.1 Roadmapping and Backlog Grooming
- 8.2 Cross-Functional Sprint Cycles (2–4 Week Sprints)
- 8.3 Stakeholder Roles: Tech, Design, and the Data Team

9. Phase 4: Product Distribution (Growth)

- 9.1 Go-to-Market (GTM) Strategy and Stakeholders
- 9.2 Viral Growth Loops and Feedback Mechanisms

Part IV: The AIPM Framework for Success

10. De-Risking the Product: The Four Risks Framework

- 10.1 Valuable: Solving Real Needs and Identifying Personas
- 10.2 Usable: Interface Optimization and Friction Reduction
- 10.3 Feasible: Technical Complexity and Scalability
- 10.4 Viable: Economic Sustainability and Strategic Fit

11. Case Study Deep Dives

11.1 Granola/Fireflies: AI Meeting Assistants

- 11.1.1 Engineering Orchestration: Transcription to System Prompts

11.2 Perplexity: Accelerating the Research Phase

11.3 Notion AI and Stripe: Personalizing the LLM Experience

12. The 7 Building Blocks of AI Products

12.1 The Model Layer: Foundational, Fine-Tuned, or Custom

12.2 The Data Layer: Cleaning, Labeling, and Storage

12.3 The Infrastructure Layer: Hosting, Monitoring, and Gpus

12.4 The UI Layer: Latency Management and Visual Guidance

12.5 Trust, Security, and Feedback Loops

12.6 Prompt Engineering Systems

12.7 The Business Layer: OKRs, ROI, and Monetization

Part V: Implementation—Context and Retrieval

13. Context Engineering Strategies

13.1 The Limitation of Context Windows: Accuracy vs. Input Size

13.2 Defining the "Art and Science" of Contextual Preparation

14. RAG: Retrieval Augmented Generation

14.1 The Step-by-Step RAG Pipeline

- 14.1.1 Data Chunking: Small and Meaningful Segments
- 14.1.2 Vector Embeddings: The Numerical Identity
- 14.1.3 Vector Storage: Pinecone and Database Selection
- 14.1.4 Augmentation and Prompt Composition

14.2 Advanced RAG: Hierarchical Systems for Summarization

15. Fine-Tuning: Deep Model Customization

15.1 Domain-Specific Behavior vs. General Knowledge

15.2 Full Fine-Tuning vs. Parameter Efficient (PEFT/LoRA)

15.3 Decision Matrix: Prompting vs. RAG vs. Fine-Tuning

16. Transfer Learning and SLMs

16.1 The Power of "Sword vs. Needle": Specializing Small Models

16.2 Synthetic Data Generation from Frontier Models

Part VI: Prompt Engineering for Professionals

17. The 20/80 Rule of Prompt Engineering

- 17.1 Role Definition (Persona Adoption)
- 17.2 Specific Output Expectations (JSON, Tables, Formats)
- 17.3 Multi-Step Thinking (Chain-of-Thought)
- 17.4 Zero-Shot vs. Few-Shot Prompting (Providing Examples)
- 17.5 Explicit Constraints and Iterative Refinement

18. The Prompt Engineer's Toolkit

- 18.1 Building a Prompt Directory and Manager
- 18.2 Anthropic's Prompt Creator and Wild Card Strategies

Part VII: Advanced AI Capabilities—Agents and MCP

19. Building AI Agents

- 19.1 Intelligence + Action + Autonomy
- 19.2 Workflow vs. Autonomous Agents
- 19.3 The Agent Architecture: Reasoning, Memory, and Tools
- 19.4 Operational Hazards: Why 95% of AI Initiatives Fail

20. The Model Context Protocol (MCP)

- 20.1 Solving the API Nightmare: A Standardized Interaction Layer
- 20.2 The Three Components: Host, Client, and Server
- 20.3 Case Study: Zomato's MCP Implemen

21. Agent Orchestration Case Studies

- 21.1 Lovable, Bolt, and V0: Vibe Coding Agents
- 21.2 Back-End Orchestration: The Step-by-Step Loop

Part VIII: Quality Control and Evaluation

22. Taming the Stochastic Beast: Evaluations

- 22.1 Identifying Challenges: Hallucinations and Indeterminism
- 22.2 The Evaluation Decision Matrix: Model Choice vs. Latency

23. LLM-as-a-Judge

- 23.1 High-Level Quality Checking: Syntactic vs. Factual Correctness
- 23.2 Case Study: Evaluating an AI-First Job Website
 - 23.2.1 Drafting the Evaluation Prompt Checklist

Part IX: The Professional Roadmap

24. Transitioning into AI Product Management

- 24.1 Converting Information to Knowledge Through Action
- 24.2 Seven Steps to Proving Competency
 - 24.2.1 Reverse Engineering Top AI Products
 - 24.2.2 Solving Workflow Gaps with Personal AI Tools
 - 24.2.3 Articulating Thought Leadership: Writing to Learn
 - 24.2.4 Tracking and Improving Emerging AI Startups

25. Portfolio Development and Networking

- 25.1 Building a "Proof of Work" Portfolio
- 25.2 Identifying Opportunities in the Y Combinator Ecosystem

26. The Job Search Strategy: The Three Bucket Approach

- 26.1 Bucket 1 (80%): High Volume and Relevant Resumes
- 24.2 Bucket 2 (15%): Personalized Cover Letters and Portfolios
- 24.3 Bucket 3 (5%): Value-Add Research and "First 6 Months" Plans

Part X: Resource Guides

- 27. The AIPM Tool Discovery Kit
- 28. Recommended Literature (The Mom's Test, Marty Cagan)
- 29. Glossary of AI Terms

Comment HELLOPM 📌
for complete guide