

1) Business Context

Last updated by | Mennell, Jacob | 19 May 2026 at 12:10 BST

Business Context

Purpose

Explain the claims triage problem, the intended business outcomes, and the user groups this solution supports.

Business Problem

Claims handling teams often work across large volumes of documents, fragmented claim records, and time-sensitive operational decisions. Important information is distributed across correspondence, reports, structured claim fields, and supporting evidence, which makes early triage and review difficult to do quickly and consistently.

In a traditional workflow, adjusters and operations teams may need to:

- open and review multiple documents to locate key facts
- compare narrative evidence against structured claim data
- decide which claims need deeper review or specialist attention
- repeat similar extraction and checking steps across many claims

This creates operational friction. Review effort can become highly manual, output quality can vary by reviewer, and teams can spend substantial time locating evidence before they can make or validate a decision.

Why Augmented Triage Is Needed

The solution is designed to improve this early-stage claims handling process by combining conventional application workflows with AI-assisted document understanding, evidence retrieval, extraction, and checking.

The aim is not to replace human decision-making. The aim is to reduce the amount of manual effort needed to find relevant claim evidence, assemble context, and surface potentially useful outputs for review.

An augmented triage approach is useful when the business needs to:

- process document-heavy claims work at greater speed
- improve consistency across repeated review tasks
- surface evidence-backed outputs rather than relying only on manual document reading
- support users with different roles, from operational reviewers to administrators and model owners

Intended Business Outcomes

At a business level, the platform is intended to support several outcomes:

- faster movement from document receipt to actionable claim insight
- more consistent extraction of relevant information from supporting evidence
- earlier identification of issues, gaps, or follow-up actions through automated checks
- better use of specialist review time by reducing low-value manual evidence gathering
- improved operational visibility into workflow execution and processing status

These outcomes matter because claims teams typically need both speed and traceability. Faster handling alone is not enough if users cannot understand where outputs came from or how they should be reviewed.

User Groups Supported

The repository indicates several user groups and operational personas:

- claims handlers and reviewers who need to inspect claim information, documents, extracted outputs, and checks

The frontend structure supports this reading. It includes claim-focused views, extracted-information pages, check-oriented pages, model-specific claim views, and an administrative area with protected access.

How The Solution Supports Those Users

The solution supports these groups through a combination of API, workflow, and user interface capabilities:

- a claim-centric web experience for navigating claims and related outputs
- API endpoints for claim retrieval, orchestration, model access, and computation
- asynchronous processing for workflow execution rather than forcing long-running work into synchronous user interactions
- retrieval over ingested claim evidence so AI-assisted extraction and checking can use grounded document context

This means the platform is positioned as an operational claims tool rather than a standalone AI demonstration. The AI components sit within a broader business workflow that includes ingestion, access control, orchestration, review, and support functions.

Business Framing

From an external perspective, Claims Augmented Triage can be described as a claims-processing platform that helps teams work more effectively with complex claim evidence. It combines document ingestion, retrieval, AI-assisted analysis, and workflow execution so that users can move from raw documents to structured, reviewable outputs more efficiently.

The value proposition is strongest in document-heavy and operationally repetitive scenarios where users need help locating relevant evidence, producing consistent outputs, and prioritizing human attention where it adds the most value.